

(<https://colab.research.google.com/github/Python-Charmer/Final-Project-Team-Python-Charmer/blob/master/Phase%203/Code/FinalProject.ipynb>)

Python Final Project - Team Python Charmers

```
In [ ]: # Loading Packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
```

```
In [ ]: # Loading Data From Source.
def load_data():
    url = r'https://raw.githubusercontent.com/Python-Charmer/Final-Project-Team-Python-Charmer/master/Phase1/Data/BreastCancerWisconsin.csv'
    df = pd.read_csv(url)
    names = ['Scn', 'A2', 'A3', 'A4', 'A5', 'A6', 'A7', 'A8', 'A9', 'A10', 'Class']
    df.columns = names
    return df
```

```
In [ ]: # Understanding Missing Values
def clean_missing(df):
    df['A7'] = df['A7'].replace('?', np.NaN)
    df['A7'] = pd.to_numeric(df['A7'])
    print("Below are how many missing values for each column\n")
    print(df.isnull().sum())
    print("\nCleaning missing values with column means\n")
    df = df.fillna(round(df.mean(skipna = True), 2))
    print(df.isnull().sum())
    return df
```

```
In [ ]: # Calculating Summary Metrics
def sum_metrics(df):
    print("\n Below are the summary metrics of the data \n" + str(df.describe()))
    print("\n\nThere are " + str(df.shape[0]) + " rows and " + str(df.shape[1]) + " Columns in this data frame")
    print("\n\nThere are " + str(len(df['Scn'].unique())) + " unique scn values in the dataset.\n")
    print("Below are the duplicate rows in the dataset.\n")
    print(str(df.loc[df.duplicated(), :]) + "\n")
```

```
In [ ]: # Plotting graphs
def plot_graphs(df):
    print("\nBelow are the histograms of A2:A10 \n")
    df.iloc[:, 1:10].hist(bins = 8, color="blue", grid="False", alpha = .5, figsize=(12,6))
    plt.tight_layout(rect=(0,0,1.2,1.2))
    plt.show()
    df['Class'].value_counts().plot.bar().set_title("Class Variable: 2 = Benign 4 = Malignant")
    df.plot.scatter(x='A3', y='A4').set_title("Scatter of A3 & A4 90% corr")
```

```
In [ ]: # We are getting centers for K = 4 clusters
def get_mids(X):
    clss = KMeans(n_clusters = 4)
    clss.fit(X)
    cent = clss.cluster_centers_
    print("\n Below are the centers of K = 4 clusters \n")
    print(pd.DataFrame(cent ,columns = X.columns))
```

```
In [ ]: # We are plotting inertia plot to find optimal K
def find_optimal_K(X):
    print("\n Below is the inertia chart \n")
    inertia = []
    k = []
    for i in range(1,15):
        clss = KMeans(n_clusters = i)
        clss.fit(X)
        iner = clss.inertia_
        k.append(i)
        inertia.append(iner)
    res = pd.concat([pd.DataFrame(k), pd.DataFrame(inertia)],axis = 1)
    res.columns = ['K','Inertia']
    ax = res.plot("K",marker='o', linestyle='dashed', title = "Optimal K = 2" )
    ax.set_xlabel("Number of Clusters")
    ax.set_ylabel("Inertia")
```

```
In [ ]: # Plotting SD plot to understand the data variance
def sd_plot(X):
    dt = pd.DataFrame(X.std()).sort_values(by = 0, ascending = False)
    dt.reset_index()
    fig, ay = plt.subplots()
    x_val = dt.index
    y_val = dt[0].values
    ay.bar(x = x_val, height = y_val)
    ay.set_xlabel("Features")
    ay.set_ylabel("Standard Deviation")
    ay.set_title("Standard Deviation Plot")

# Plotting Box plot to understand the data variance
def var_plot(df):
    # Box plot showing variation of the columns A2:A10
    data = []
    for i in range(1, 10):
        data.append(df.iloc[:, i])

    # Multiple box plots on one Axes
    fig, ax = plt.subplots()
    plt.title("Boxplot showing Variation of Features")
    plt.xlabel("Columns A2 thru A10")
    plt.ylabel("Values")
    ax.boxplot(data, 0,showbox=True,showmeans=True)
    top = 12
    bottom = -2
    ax.set_ylim(bottom, top)
    ax.set_xticklabels(df.iloc[:,1:-1].columns, rotation=45, fontsize=8)
    plt.show()
```

```
In [ ]: #Getting centers of optimal K = 2
def get_centers(X):
    print("\n Below are the centers of K = 2 clusters \n \n")
    mdl = make_pipeline(StandardScaler(), KMeans(n_clusters = 2, n_init=20))
    mdl.fit(X)
    centers = pd.DataFrame(mdl.named_steps['kmeans'].cluster_centers_)
    centers.columns = X.columns
    print(centers)
```

```
In [ ]: # Cross tabulating the cluster labels with "Class"
def lables(i,df):
    print("\nBelow are the predicted labels with k = " + str(i) + "\n")
    if i == 4:
        mdl = KMeans(n_clusters = i)
    else:
        mdl = make_pipeline(StandardScaler(), KMeans(n_clusters = i, n_init=20))
    labels = mdl.fit_predict(df.iloc[:,1:-1])
    ctf = pd.DataFrame({'labels': labels, 'Class': df["Class"]})
    print(pd.crosstab(ctf['labels'], ctf['Class']))
```

```
In [11]: # Main Function Phase 1
df = load_data()
df = clean_missing(df)
sum_metrics(df)
plot_graphs(df)
print("The columns that need standardization are: A7,A3,& A9 because they have the highest amount of variance compared to other factors.")
```

Below are how many missing values for each column

```
Scn      0
A2       0
A3       0
A4       0
A5       0
A6       0
A7      16
A8       0
A9       0
A10      0
Class    0
dtype: int64
```

Cleaning missing values with column means

```
Scn      0
A2       0
A3       0
A4       0
A5       0
A6       0
A7       0
A8       0
A9       0
A10      0
Class    0
dtype: int64
```

Below are the summary metrics of the data

	Scn	A2	A3	A4	A5 \
count	6.990000e+02	699.000000	699.000000	699.000000	699.000000
mean	1.071704e+06	4.417740	3.134478	3.207439	2.806867
std	6.170957e+05	2.815741	3.051459	2.971913	2.855379
min	6.163400e+04	1.000000	1.000000	1.000000	1.000000
25%	8.706885e+05	2.000000	1.000000	1.000000	1.000000
50%	1.171710e+06	4.000000	1.000000	1.000000	1.000000
75%	1.238298e+06	6.000000	5.000000	5.000000	4.000000
max	1.345435e+07	10.000000	10.000000	10.000000	10.000000

	A6	A7	A8	A9	A10	Class
count	699.000000	699.000000	699.000000	699.000000	699.000000	699.000000
mean	3.216023	3.544549	3.437768	2.866953	1.589413	2.689557
std	2.214300	3.601852	2.438364	3.053634	1.715078	0.951273
min	1.000000	1.000000	1.000000	1.000000	1.000000	2.000000
25%	2.000000	1.000000	2.000000	1.000000	1.000000	2.000000
50%	2.000000	1.000000	3.000000	1.000000	1.000000	2.000000
75%	4.000000	5.000000	5.000000	4.000000	1.000000	4.000000
max	10.000000	10.000000	10.000000	10.000000	10.000000	4.000000

There are 699 rows and 11 Columns in this data frame


There are 645 unique scn values in the dataset.

Below are the duplicate rows in the dataset.


	Scn	A2	A3	A4	A5	A6	A7	A8	A9	A10	Class
208	1218860	1	1	1	1	1	1.0	3	1	1	2
253	1100524	6	10	10	2	8	10.0	7	3	3	4
254	1116116	9	10	10	1	10	8.0	3	3	1	4
258	1198641	3	1	1	1	2	1.0	3	1	1	2

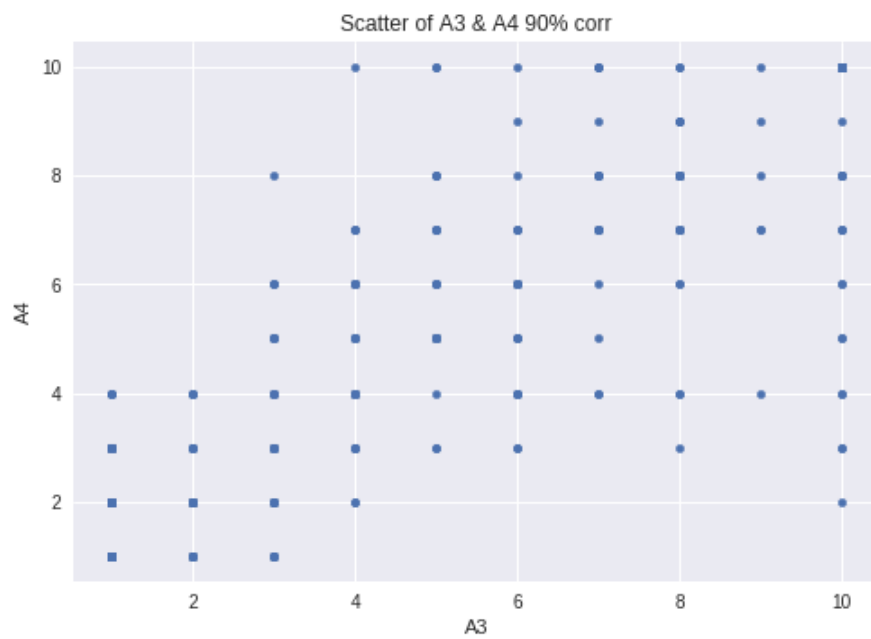
272	320675	3	3	5	2	3	10.0	7	1	1	4
338	704097	1	1	1	1	1	1.0	2	1	1	2
561	1321942	5	1	1	1	2	1.0	3	1	1	2
684	466906	1	1	1	1	2	1.0	1	1	1	2

Below are the histograms of A2:A10



The columns that need standardization are: A7,A3,& A9 because they have the highest amount of variance compared to other factors.





```
In [12]: #Main Functions Phase 2
X = df.drop(['Scn', 'Class'], axis = 1)
y = df['Class']
get_mids(X)
lables(4,df)
find_optimal_K(X)
sd_plot(X)
var_plot(df)
print('\n Based on the Box and SD plot above we can see features A7,A9 has the most variations.\n')
get_centers(X)
lables(2,df)
```


Below are the centers of K = 4 clusters


	A2	A3	A4	A5	A6	A7	A8	\
0	7.204082	4.846939	5.010204	4.816327	4.071429	9.158571	5.224490	
1	2.984716	1.266376	1.386463	1.312227	2.054585	1.352576	2.080786	
2	6.721519	8.367089	8.405063	7.810127	6.734177	9.227848	7.367089	
3	7.562500	7.421875	7.062500	4.250000	5.875000	3.619063	5.562500	

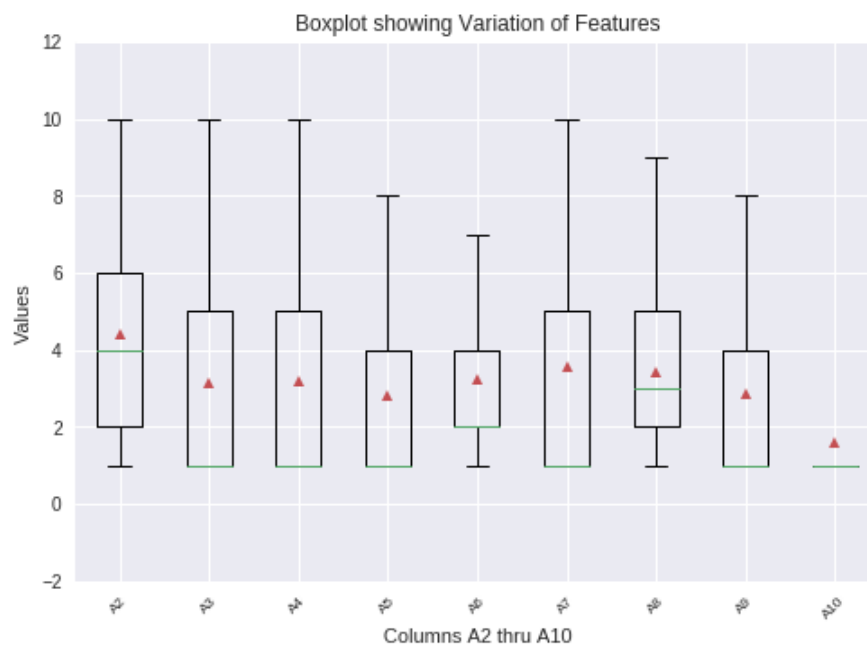
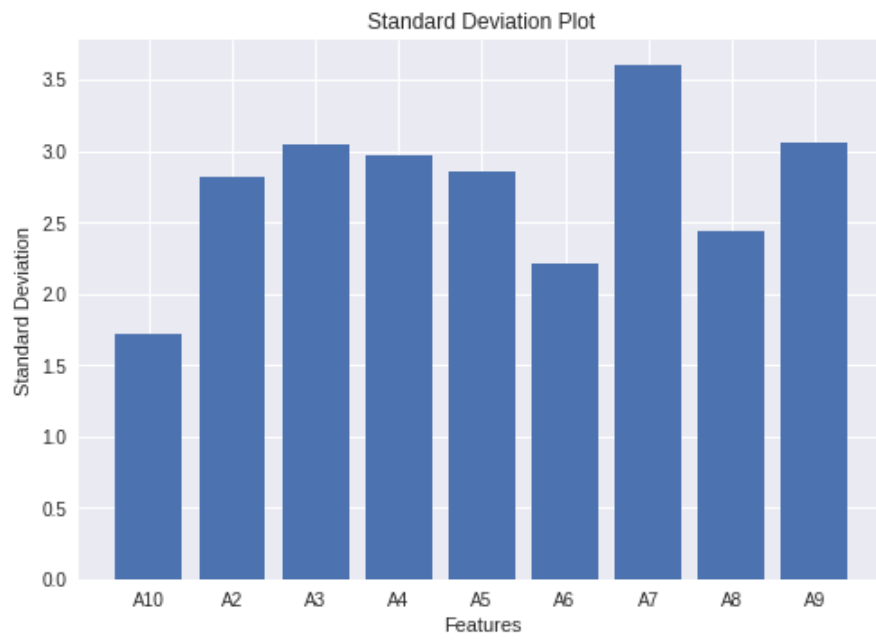
	A9	A10
0	3.795918	1.642857
1	1.213974	1.102620
2	7.822785	3.822785
3	7.156250	2.234375

Below are the predicted labels with k = 4

Class	2	4
labels		
0	7	64
1	444	10
2	7	87
3	0	80

Below is the inertia chart





Based on the Box and SD plot above we can see features A7,A9 has the most variations.

Below are the centers of K = 2 clusters

	A2	A3	A4	A5	A6	A7	A8	\
0	-0.496223	-0.60690	-0.602092	-0.514917	-0.509713	-0.580604	-0.547702	
1	0.986083	1.20602	1.196465	1.023233	1.012892	1.153765	1.088383	

	A9	A10
0	-0.530778	-0.303758
1	1.054751	0.603622

Below are the predicted labels with k = 2

Class	2	4
labels		
0	446	19
1	12	222

```

In [25]: #Main Phase 3
mdl = make_pipeline(StandardScaler(), KMeans(n_clusters = 2, n_init=20, max_iter = 500
))
labels = mdl.fit_predict(X)
df['Predicted'] = labels

for x in range(df.shape[0]):
    if df.iloc[x,11] == 0:
        df.iloc[x,11] = 2
    else:
        df.iloc[x,11] = 4

print("\nBelow are the first 15 rows of the dataframe \n")
print(df.head(15))

print("\nBelow are the observtions where the predicted did not match the class \n")

print(df[df['Class'] != df['Predicted']])

def error_rate(predicted,actual):
    tab = pd.crosstab(actual,predicted)
    error2 = tab.iloc[0,1]
    total2 = tab.iloc[0,0] + tab.iloc[1,0]

    error4 = tab.iloc[1,0]
    total4 = tab.iloc[0,1] + tab.iloc[1,1]

    B = str(round(error2/total2,4)*100) + "%"
    M = str(round(error4/total4,4)*100) + "%"
    tot_error = str(round((error2 + error4)/(total2 + total4),4)*100) + "%"

    print("\nThe error rate for beningn cells is " + str(B) + "\n")
    print("The error rate for malignant cells is " +str(M) + "\n")
    print("The total error rate is " +str(tot_error) + "\n")

error_rate(df['Predicted'], df['Class'])

```

Below are the first 15 rows of the dataframe

	Scn	A2	A3	A4	A5	A6	A7	A8	A9	A10	Class	Predicted
0	1000025	5	1	1	1	2	1.0	3	1	1	2	2
1	1002945	5	4	4	5	7	10.0	3	2	1	2	4
2	1015425	3	1	1	1	2	2.0	3	1	1	2	2
3	1016277	6	8	8	1	3	4.0	3	7	1	2	4
4	1017023	4	1	1	3	2	1.0	3	1	1	2	2
5	1017122	8	10	10	8	7	10.0	9	7	1	4	4
6	1018099	1	1	1	1	2	10.0	3	1	1	2	2
7	1018561	2	1	2	1	2	1.0	3	1	1	2	2
8	1033078	2	1	1	1	2	1.0	1	1	5	2	2
9	1033078	4	2	1	1	2	1.0	2	1	1	2	2
10	1035283	1	1	1	1	1	1.0	3	1	1	2	2
11	1036172	2	1	1	1	2	1.0	2	1	1	2	2
12	1041801	5	3	3	3	2	3.0	4	4	1	4	2
13	1043999	1	1	1	1	2	3.0	3	1	1	2	2
14	1044572	8	7	5	10	7	9.0	5	5	4	4	4

Below are the observtions where the predicted did not match the class

	Scn	A2	A3	A4	A5	A6	A7	A8	A9	A10	Class	Predicted
1	1002945	5	4	4	5	7	10.00	3	2	1	2	4
3	1016277	6	8	8	1	3	4.00	3	7	1	2	4
12	1041801	5	3	3	3	2	3.00	4	4	1	4	2
25	1065726	5	2	3	4	2	7.00	3	6	1	4	2
40	1096800	6	6	6	9	6	3.54	7	8	1	2	4
51	1108449	5	3	3	4	2	4.00	3	4	1	4	2
57	1113038	8	2	4	1	5	1.00	5	4	4	4	2
58	1113483	5	2	3	1	6	10.00	5	1	1	4	2
59	1113906	9	5	5	2	2	2.00	5	1	1	4	2
63	1116132	6	3	4	1	5	2.00	3	9	1	4	2
101	1167439	2	3	4	4	2	5.00	2	5	1	4	2
103	1168359	8	2	3	1	6	3.00	7	1	1	4	2
146	1185609	3	4	5	2	6	8.00	4	1	1	4	2
179	1202812	5	3	3	3	6	10.00	3	1	1	4	2
196	1213375	8	4	4	5	4	7.00	7	8	2	2	4
222	1226012	4	1	1	3	1	5.00	2	1	1	4	2
247	145447	8	4	4	1	2	9.00	3	3	1	4	2
252	1017023	6	3	3	5	3	10.00	3	5	3	2	4
259	242970	5	7	7	1	5	8.00	3	4	1	2	4
273	428903	7	2	4	1	3	4.00	3	3	1	4	2
296	616240	5	3	4	3	4	5.00	4	7	1	2	4
315	704168	4	6	5	6	7	3.54	4	9	1	2	4
319	721482	4	4	4	4	6	5.00	7	3	1	2	4
326	752904	10	1	1	1	2	10.00	5	4	1	4	2
348	832226	3	4	4	10	5	1.00	3	3	1	4	2
352	846832	3	4	5	3	7	3.00	4	6	1	2	4
356	859164	5	3	3	1	3	3.00	3	3	3	4	2
434	1293439	6	9	7	5	5	8.00	4	2	1	2	4
455	1246562	10	2	2	1	2	6.00	1	1	2	4	2
489	1084139	6	3	2	1	3	4.00	4	1	1	4	2
657	1333877	5	4	5	1	8	1.00	3	6	1	2	4

The error rate for beningn cells is 2.58%

The error rate for malignant cells is 8.12%

The total error rate is 4.43%

In []: